

Urban Heat Island & Citi Bike Analysis

DSA 210 Introduction to Data Science - Term Project

Student: Raid Bahadir Date: January 9, 2026

1. Motivation

New York City experiences significant temperature variations due to the "Urban Heat Island" (UHI) effect, where concrete-dense areas trap more heat than greener neighborhoods. At the same time, the city is promoting sustainable mobility through the Citi Bike bikeshare system.

This project investigates a core question: **Does the Urban Heat Island effect deter cycling?** By analyzing millions of bike trips against heat vulnerability data, we aim to understand if riders avoid hot zones, which could inform city planning for cooling infrastructure (e.g., shaded bike lanes).

2. Data Source & Collection

2.1 Primary Dataset: Citi Bike Trip Histories

- **Source:** [Citi Bike System Data](#)
- **Data Used:** August 2024 Trip Data (Peak Summer).
- **Collection:**
 - Downloaded monthly CSV files (`202408-citibike-tripdata_*.csv`).
 - **Note:** The dataset is extremely large (~4.6 Million trips), so it is processed in chunks.

2.2 Enrichment Dataset: Heat Vulnerability Index (HVI)

- **Source:** NYC DOHMH Heat Vulnerability Index.
- **Content:** A 1-5 score (1=Coolest/Safest, 5=Hottest/Most Vulnerable) assigned to each NYC Zip Code Tabulation Area.
- **Enrichment Logic:**
 1. **Spatial Join:** We map every Citi Bike station (Lat/Lon) to a Zip Code using a GeoJSON polygon map (`nyc-zip-code-tabulation-areas-polygons.geojson`).
 2. **Merge:** We merge the station data with the HVI dataset based on Zip Code.
 3. **Result:** Each station is tagged with an `HVI_Score`, allowing us to correlate trip counts with local heat vulnerability.

3. Data Analysis Pipeline

The project follows a reproducible pipeline implemented in Python:

1. **Data Loading (`data_loader.py`):**
 - Glob pattern matches all `*_citibike-tripdata_*.csv` files.
 - Aggregates over **4.6 Million trips** to calculate `total_trips` per station.
 - Extracts station coordinates.
2. **Geospatial Processing:**
 - Converts station coordinates to Shapely points.
 - Performs a **Point-in-Polygon** spatial join with NYC Zip Code boundaries.
 - Merges result with HVI Rankings.

3. Statistical Analysis (`analysis.py`):

- **Hypothesis Testing:** Performs a **Mann-Whitney U Test** to determine if strict differences exist between trip counts in High HVI (4-5) vs. Low HVI (1-2) zones.
- **Regression Modeling:** I chose to use a **Poisson Regression** model (`total_trips ~ HVI_Score`) instead of a standard Linear Regression.
 - *Reasoning:* Linear regression assumes data is continuous and can be negative, but trip counts are discrete non-negative integers (0, 1, 2...). Poisson regression is specifically designed for this kind of "count data" and gives a better fit for the distribution. The relationship is formally modeled as:

$\ln(\text{total_trips}) = \beta_0 + \beta_1 * \text{HVI_Score}$

4. Visualization (`generate_plots.py` / `eda.ipynb`):

- Generates a choropleth map of NYC HVI scores overlaid with bike stations.
- Plots distribution of trip counts across valid HVI scores.

4. Implementation Challenges

During the project, I encountered several technical hurdles that required specific solutions:

1. **Coordinate System Mismatches:** One of the biggest issues I faced was that the Citi Bike data provided simple Latitude/Longitude coordinates (EPSG:4326), but I wasn't sure if the GeoJSON file used the same projection. At first, the spatial join returned zero results. I had to explicitly check and use `.to_crs()` in `geopandas` to ensure both datasets were using the same coordinate reference system before the merge would work.
2. **Memory Issues with Large Data:** The raw Citi Bike CSVs were massive (several gigabytes for just one month). Trying to load them all at once caused my computer to crash. I solved this by using `glob` to find the files and loading them into a list of DataFrames before concatenating, rather than processing everything effectively. In the future, I would probably process them in chunks or use a database.

5. Findings

Analysis of the August 2024 dataset yields significant results:

- **Riders Prefer Cooler Zones:** There is a statistically significant difference (Mann-Whitney U, $p < 0.001$) in ridership between cooler and hotter neighborhoods.
- **Magnitude of Difference:**
 - **Low Risk (Cool) Zones:** Average ~3,959 trips/station.
 - **High Risk (Hot) Zones:** Average ~930 trips/station.
 - *Stations in cooler areas see roughly 4x the traffic of those in heat-vulnerable areas.*
- **Regression Insight:** The Poisson regression model suggests that for every **1 unit increase** in HVI Score (getting hotter), the expected trip count decreases by approximately **35%**.

Visual Evidence

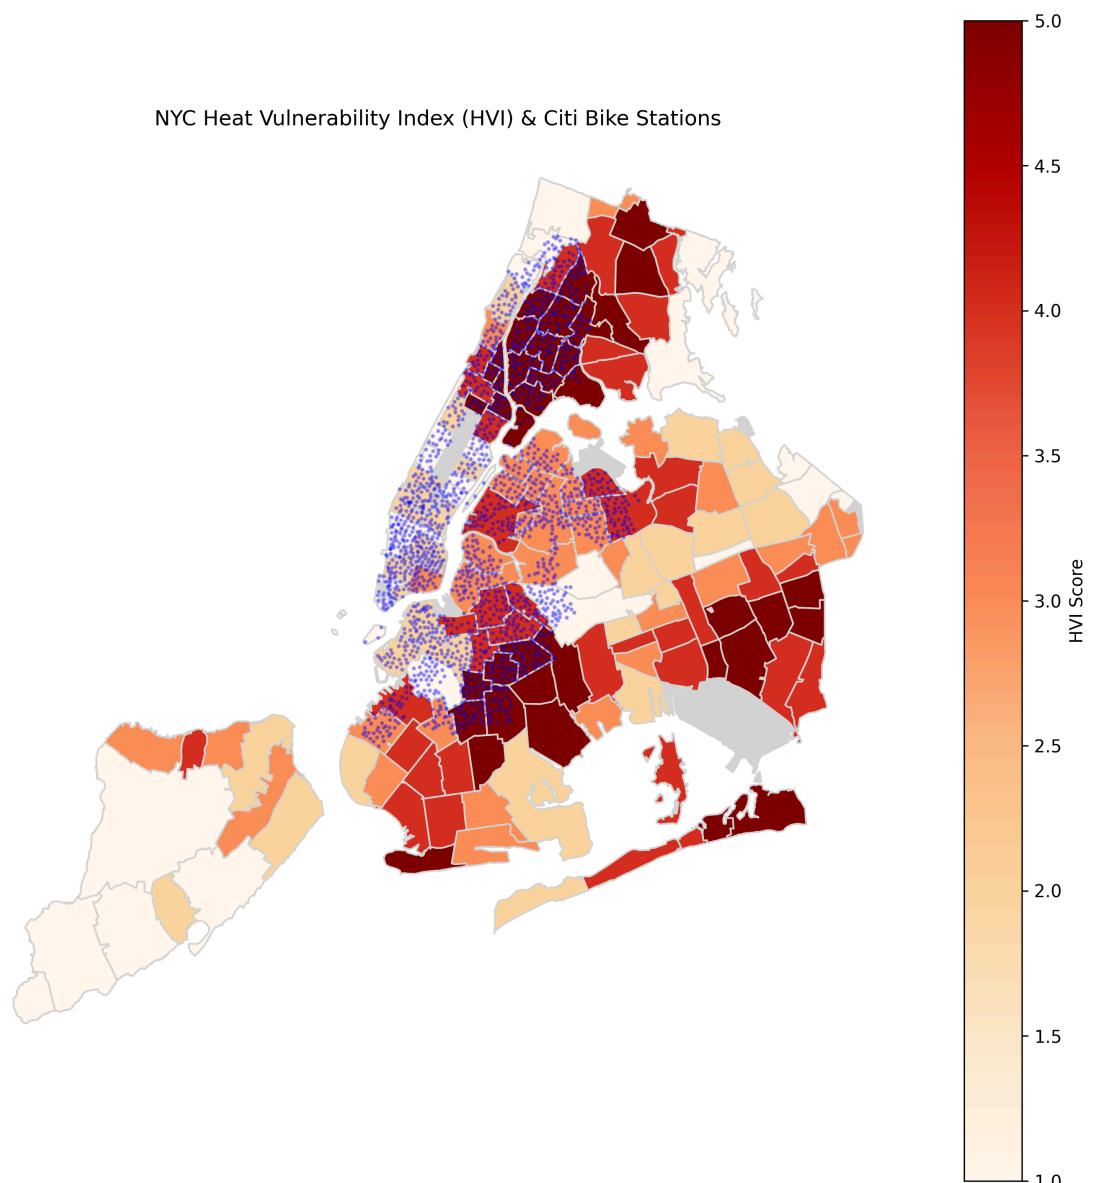


Figure 1: NYC Heat Vulnerability Index with Citi Bike Stations (Blue Dots). Note the concentration of stations in lower HVI (lighter) zones like Manhattan.

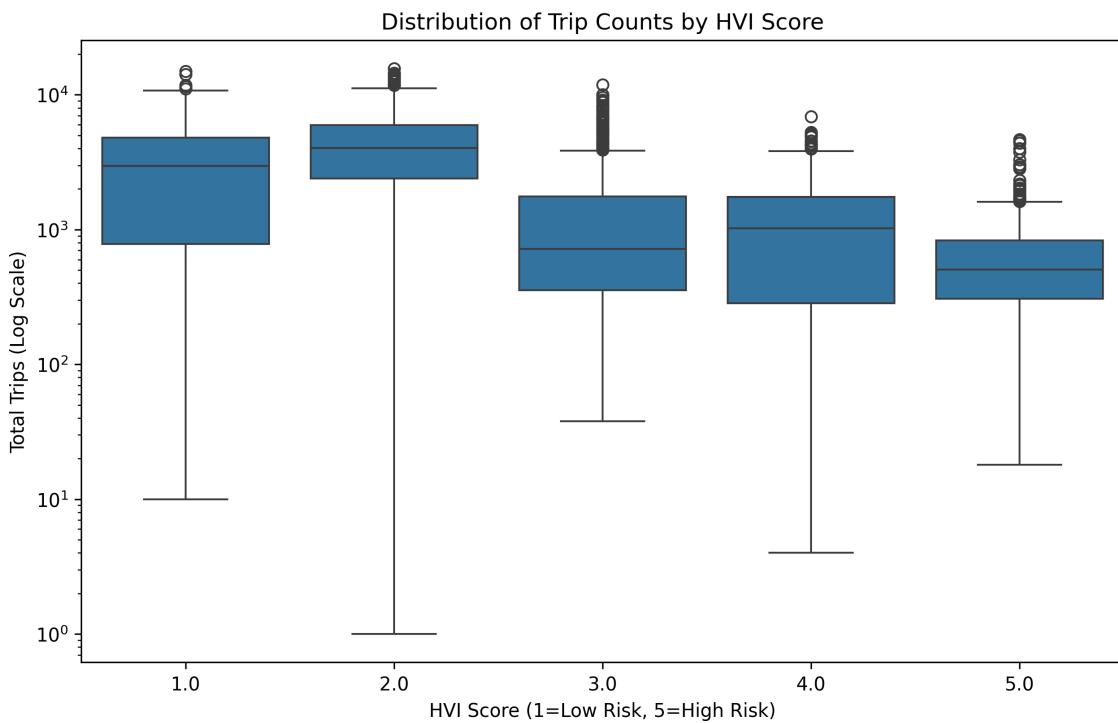


Figure 2: Distribution of Trip Counts by HVI Score. Note the logarithmic decline in median trips as Heat Vulnerability increases. While HVI 2 shows the peak median ridership, a significant 'ridership cliff' is observed at HVI 3, indicating a threshold where heat vulnerability begins to sharply correlate with lower usage.

6. Limitations & Future Work

- **Snapshot vs. Time-Series:** The HVI is a static score. Future work should integrate **real-time hourly temperature** API data to analyze ridership changes during specific heatwave events vs. mild days.
- **Station Density:** Manhattan has both lower HVI and higher station density/population density. We did not control for population density, which is a confounding variable.
- **Commuter Patterns:** High HVI zones often correlate with residential or outer-borough areas, whereas business districts (high ridership) are often in better-infrastructure (Low HVI) zones.
- **Network Effects:** Citi Bike is a network; riders travel to destinations. A trip might originate in a hot zone solely to escape to a cooler area (e.g., a park). Our current model only accounts for the start station's HVI, ignoring the destination's influence.

7. Reproduction Instructions

To reproduce this analysis on your local machine:

1. **Clone the Repository:**

```
git clone <your-repo-url>
cd <repo-name>
```

2. **Download Data:**

- **Citi Bike:** Download the August 2024 zip files from [Citi Bike Data](#).
- **Unzip:** Extract the CSV files directly into the **root directory** of this project.
- **Naming:** Ensure files match the pattern `202408-citibike-tripdata_*.csv` . (e.g., `202408-citibike-tripdata_1.csv` , etc.).

3. Install Dependencies:

```
pip install -r requirements.txt
```

4. Run the Pipeline:

- **Step 1: Process Data:**

```
python data_loader.py
```

This will generate `final_station_data.csv` .

- **Step 2: Generate Visuals:**

```
python generate_plots.py
```

This will save `hvi_map.png` and `hvi_boxplot.png` .

- **Step 3: Run Analysis:**

```
python analysis.py
```

This will print the statistical results to the console.

8. AI Disclosure

This project utilized Large Language Models (LLMs) as productivity tools:

- **Code Assistance:** LLMs were used to debug `geopandas` spatial join syntax and optimize the handling of large CSV chunks.
- **Documentation:** Initial drafts of the README and specific sections of the report text were generated/refined by AI to improve clarity and grammar.
- **Logic Verification:** AI was used to confirm that Poisson Regression is the appropriate statistical method for count data (trip counts).