

# Urban Heat Island & Citi Bike Analysis

DSA 210 Introduction to Data Science - Term Project

Student: Raid Bahadir Date: January 9, 2026

---

## 1. Motivation

New York City experiences significant temperature variations due to the "Urban Heat Island" (UHI) effect, where concrete-dense areas trap more heat than greener neighborhoods. At the same time, the city is promoting sustainable mobility through the Citi Bike bikeshare system.

This project investigates a core question: **Does the Urban Heat Island effect deter cycling?** By analyzing millions of bike trips against heat vulnerability data, we aim to understand if riders avoid hot zones, which could inform city planning for cooling infrastructure (e.g., shaded bike lanes).

## 2. Data Source & Collection

### 2.1 Primary Dataset: Citi Bike Trip Histories

- **Source:** [Citi Bike System Data](#)
- **Data Used:** August 2024 Trip Data (Peak Summer).
- **Collection:**
  - Downloaded monthly CSV files ( `202408-citibike-tripdata_*.csv` ).
  - **Note:** The dataset is extremely large (~4.6 Million trips), so it is processed in chunks.

### 2.2 Enrichment Dataset: Heat Vulnerability Index (HVI)

- **Source:** NYC DOHMH Heat Vulnerability Index.
- **Content:** A 1-5 score (1=Coolest/Safest, 5=Hottest/Most Vulnerable) assigned to each NYC Zip Code Tabulation Area.
- **Enrichment Logic:**
  1. **Spatial Join:** We map every Citi Bike station (Lat/Lon) to a Zip Code using a GeoJSON polygon map ( `nyc-zip-code-tabulation-areas-polygons.geojson` ).
  2. **Merge:** We merge the station data with the HVI dataset based on Zip Code.
  3. **Result:** Each station is tagged with an `HVI_Score`, allowing us to correlate trip counts with local heat vulnerability.

## 3. Data Analysis Pipeline

The project follows a reproducible pipeline implemented in Python:

1. **Data Loading ( `data_loader.py` ):**
  - Glob pattern matches all `*_citibike-tripdata_*.csv` files.
  - Aggregates over **4.6 Million trips** to calculate `total_trips` per station.
  - Extracts station coordinates.
2. **Geospatial Processing:**
  - Converts station coordinates to Shapely points.
  - Performs a **Point-in-Polygon** spatial join with NYC Zip Code boundaries.
  - Merges result with HVI Rankings.

### 3. Statistical Analysis ( `analysis.py` ):

- **Hypothesis Testing:** Performs a **Mann-Whitney U Test** to determine if strict differences exist between trip counts in High HVI (4-5) vs. Low HVI (1-2) zones.
- **Regression Modeling:** I chose to use a **Poisson Regression** model ( `total_trips ~ HVI_Score` ) instead of a standard Linear Regression.
  - *Reasoning:* Linear regression assumes data is continuous and can be negative, but trip counts are discrete non-negative integers (0, 1, 2...). Poisson regression is specifically designed for this kind of "count data" and gives a better fit for the distribution. The relationship is formally modeled as:

$\ln(\text{total\_trips}) = \beta_0 + \beta_1 * \text{HVI\_Score}$

### 4. Visualization ( `generate_plots.py` / `eda.ipynb` ):

- Generates a choropleth map of NYC HVI scores overlaid with bike stations.
- Plots distribution of trip counts across valid HVI scores.

## 4. Implementation Challenges

During the project, I encountered several technical hurdles that required specific solutions:

1. **Coordinate System Mismatches:** One of the biggest issues I faced was that the Citi Bike data provided simple Latitude/Longitude coordinates (EPSG:4326), but I wasn't sure if the GeoJSON file used the same projection. At first, the spatial join returned zero results. I had to explicitly check and use `.to_crs()` in `geopandas` to ensure both datasets were using the same coordinate reference system before the merge would work.
2. **Memory Issues with Large Data:** The raw Citi Bike CSVs were massive (several gigabytes for just one month). Trying to load them all at once caused my computer to crash. I solved this by using `glob` to find the files and loading them into a list of DataFrames before concatenating, rather than processing everything effectively. In the future, I would probably process them in chunks or use a database.

## 5. Findings

Analysis of the August 2024 dataset yields significant results:

- **Riders Prefer Cooler Zones:** There is a statistically significant difference (Mann-Whitney U,  $p < 0.001$ ) in ridership between cooler and hotter neighborhoods.
- **Magnitude of Difference:**
  - **Low Risk (Cool) Zones:** Average ~3,959 trips/station.
  - **High Risk (Hot) Zones:** Average ~930 trips/station.
  - *Stations in cooler areas see roughly 4x the traffic of those in heat-vulnerable areas.*
- **Regression Insight:** The Poisson regression model suggests that for every **1 unit increase** in HVI Score (getting hotter), the expected trip count decreases by approximately **35%**.

## Visual Evidence

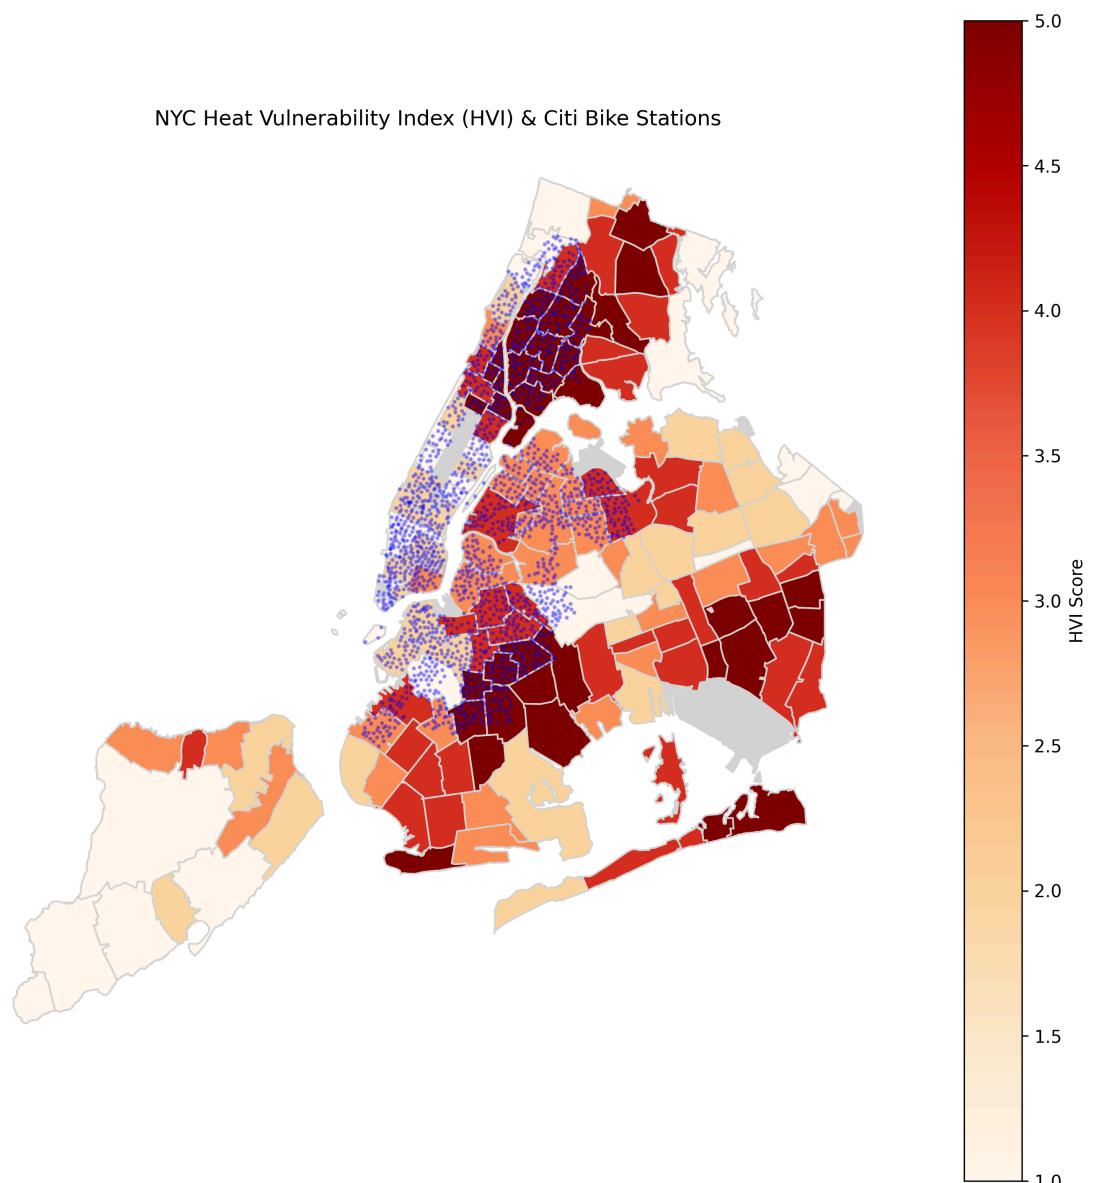
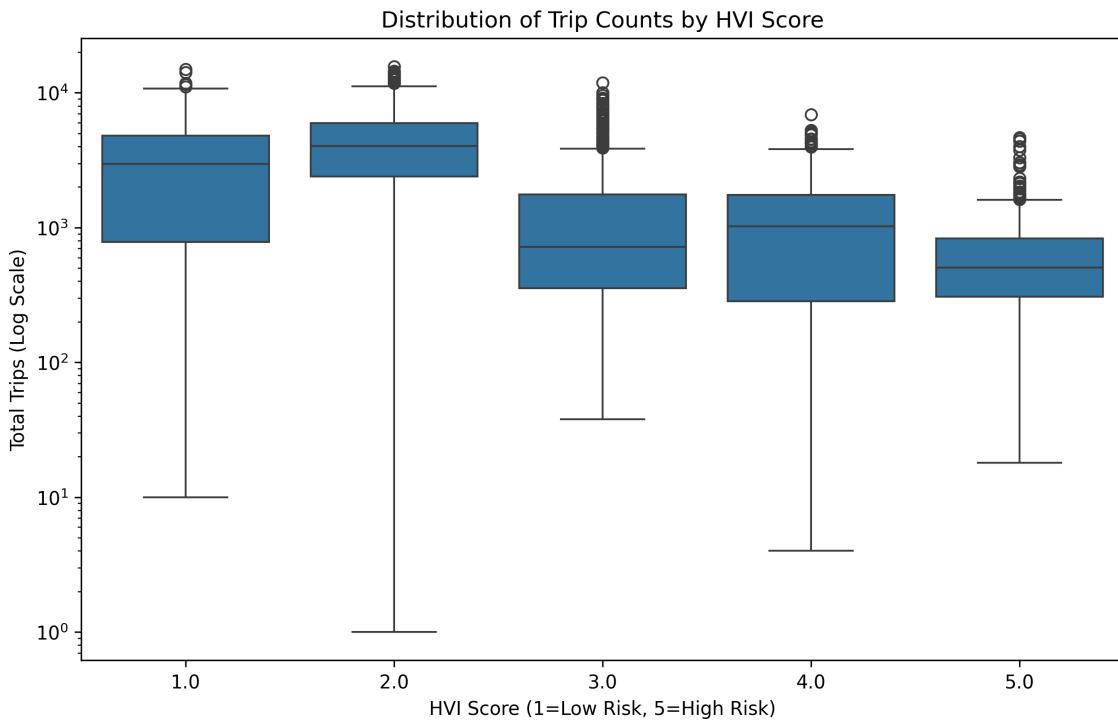


Figure 1: NYC Heat Vulnerability Index with Citi Bike Stations (Blue Dots). Note the concentration of stations in lower HVI (lighter) zones like Manhattan.



*Figure 2: Distribution of Trip Counts by HVI Score. Note the logarithmic decline in median trips as Heat Vulnerability increases. While HVI 2 shows the peak median ridership, a significant 'ridership cliff' is observed at HVI 3, indicating a threshold where heat vulnerability begins to sharply correlate with lower usage.*

## 6. Limitations & Future Work

- **Snapshot vs. Time-Series:** The HVI is a static score. Future work should integrate **real-time hourly temperature** API data to analyze ridership changes during specific heatwave events vs. mild days.
- **Station Density:** Manhattan has both lower HVI and higher station density/population density. We did not control for population density, which is a confounding variable.
- **Commuter Patterns:** High HVI zones often correlate with residential or outer-borough areas, whereas business districts (high ridership) are often in better-infrastructure (Low HVI) zones.
- **Network Effects:** Citi Bike is a network; riders travel to destinations. A trip might originate in a hot zone solely to escape to a cooler area (e.g., a park). Our current model only accounts for the start station's HVI, ignoring the destination's influence.

## 7. Reproduction Instructions

To reproduce this analysis on your local machine:

1. **Clone the Repository:**

```
git clone https://github.com/RaidBahadir/urban-heat-mobility.git
cd urban-heat-mobility
```

2. **Download Data:**

- **Citi Bike:** Download the August 2024 zip files from [Citi Bike Data](#).
- **Unzip:** Extract the CSV files directly into the **root directory** of this project.
- **Naming:** Ensure files match the pattern `202408-citibike-tripdata_*.csv` . (e.g., `202408-citibike-tripdata_1.csv` , etc.).

### 3. Install Dependencies:

```
pip install -r requirements.txt
```

### 4. Run the Pipeline:

- **Step 1: Process Data:**

```
python data_loader.py
```

*This will generate `final_station_data.csv` .*

- **Step 2: Generate Visuals:**

```
python generate_plots.py
```

*This will save `hvi_map.png` and `hvi_boxplot.png` .*

- **Step 3: Run Analysis:**

```
python analysis.py
```

*This will print the statistical results to the console.*

## 8. AI Disclosure

This project utilized Large Language Models (LLMs) as productivity tools:

- **Code Assistance:** LLMs were used to debug `geopandas` spatial join syntax and optimize the handling of large CSV chunks.
- **Documentation:** Initial drafts of the README and specific sections of the report text were generated/refined by AI to improve clarity and grammar.
- **Logic Verification:** AI was used to confirm that Poisson Regression is the appropriate statistical method for count data (trip counts).